

# Class 07: Machine Learning 1

Amy Nguyen (PID: A18148284)

## Table of contents

Background . . . . .	1
K-means clustering . . . . .	3
Hierachial clustering . . . . .	7
Principal Component Analysis (PCA) . . . . .	11
PCA of UK food data . . . . .	11
Heatmap . . . . .	16
PCA the rescue . . . . .	17
Digging deeper (variable loadings) . . . . .	23
PCA of RNA-seq data . . . . .	25


## Background

Today we will be begin our re-exploration of important machine learning methods with a focus on **clustering** and **dimensionality reduction**

To start testing these methods let's make up some sample data to cluster where we know what the answer should be.


```
hist(rnorm(3000, mean=10))
```

**Histogram of rnorm(3000, mean = 10)**



Q. Can you generate 30 numbers centered at +3 taken at random from a normal distribution?

```
tmp <- c(rnorm(30, mean=3),  
         rnorm(30, mean=-3) )  
  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



## K-means clustering

The main function in “base R” for K-means clustering is called `kmeans()`, lets try it out:

```
k <- kmeans(x, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

## Cluster means:

```

          x           y
1 -3.046488  2.839540
2  2.839540 -3.046488

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 56.28266 56.28266
```

(between\_SS / total\_SS = 90.2 %)

### Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. What component of your kmeans result object has the cluster centers?

k\$centers

	x	y
1	-3.046488	2.839540
2	2.839540	-3.046488

Q. What component of your kmeans result object has the cluster size (i.e. how many points are in each cluster)?

k\$size


[1] 30 30

Q. What component of your kmeans result object has the cluster membership vector (i.e. the main clustering result: which points are in which cluster)?

k\$cluster

Q. Plot the results of clustering (i.e. our data colored by the clustering result) along with the cluster centers.


```
plot(x, col=k$cluster)
```



```
#points(k$centers, col="blue", pch=15, cex=2)
```

Q. Can you run kmeans again and cluster into 4 clusters and plot the results just like we did above with coloring by cluster and the cluster centers shown in blue?


```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15)
```



Q. Run kmeans() again and this time produce 4 clusters and call your result object 'k4' and make a results figure like above?

```
k4 <- kmeans(x, center = 4)

plot(x, col=k4$cluster)
points(k$centers, col="blue", pch=15, cex=2)
```



**Key-point:** Kmeans will always return the clustering that we ask for (this is the “K” or “centers” in K-means)!

```
k$tot.withinss
```


```
[1] 112.5653
```

### Hierachial clustering

The main function for hierarchical clustering in base R is called `hclust()`. One of the main differences with respect to the `kmeans()` function is that you can not just pass your input data directly to `hclust()` - it needs a “distance matrix” as input. We can get this from lots of places including the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

## Cluster Dendrogram




```
d  
hclust (*, "complete")
```

We can “cut” the dendrogram or “tree” at a given height to yield our “clusters”. For this we use the function `cutree()`

```
plot(hc)  
abline(h=4.5, col="red")
```

## Cluster Dendrogram



d  
hclust (\*, "complete")


```
grps <- cutree(hc, h=10)
```

```
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. Plot our data `x` colored by the clustering result from `hclust()` and `cutree()`?

```
plot(x, col=grps)
```



```
plot(hc)
abline(h=5.6, col="red")
```

## Cluster Dendrogram



```
d  
hclust (*, "complete")
```

```
grps <- cutree(hc, h=5.6)
```

## Principal Component Analysis (PCA)

PCA is a popular dimensional reduction technique that is widely used in bioinformatics.

### PCA of UK food data

Read data on food consumption in the UK

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
x
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

7	Fresh_potatoes	720	874	566	1033
8	Fresh_Veg	253	265	171	143
9	Other_Veg	488	570	418	355
10	Processed_potatoes	198	203	220	187
11	Processed_Veg	360	365	337	334
12	Fresh_fruit	1102	1137	957	674
13	Cereals	1472	1582	1462	1494
14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

It looks like the row names are not set properly. We can fix this.

```
rownames(x) <- x[,1]
x <- x[,-1]
x
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209
Sugars		156	175	147	139
Fresh_potatoes		720	874	566	1033
Fresh_Veg		253	265	171	143
Other_Veg		488	570	418	355
Processed_potatoes		198	203	220	187
Processed_Veg		360	365	337	334
Fresh_fruit		1102	1137	957	674
Cereals		1472	1582	1462	1494
Beverages		57	73	53	47
Soft_drinks		1374	1256	1572	1506
Alcoholic_drinks		375	475	458	135
Confectionery		54	64	62	41

A better way to do this is fix the row names assignment at import time:

```
read.csv(url, row.names = 1)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions? 17 rows, and 4 columns.


```
dim(x)
```

```
[1] 17 4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances? I prefer the read.csv(url, row.names = 1)

Q3: Changing what optional argument in the above barplot() function results in the following plot? Changing beside=T to beside=F, or deleting it.

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



```
library(tidyr)


# Convert data to long format for ggplot with `pivot_longer()`
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
               names_to = "Country",
               values_to = "Consumption")
dim(x_long)
```

[1] 68 3

Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure? Change dodge to stack


```
library(ggplot2)

ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col(position = "stack") +
  theme_bw()
```



Q5: We can use the `pairs()` function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot? A point on the diagonal of a given plot represents a value where the variable on the x-axis is equal to the variable on the y-axis. Each country is being compared to/plot against itself, so the off-diagonal plots show relationships between different countries.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```




## Heatmap

We can install the **pheatmap** package with the `install.packages()` command that we used previously. Remember that we always run this in the console and not a code chunk in our quarto document.

Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set? England, Scotland, and Wales cluster together, suggesting similar food consumption patterns. Based on the color intensity of the heat map, it can be observed that N.Ireland has higher consumption of fresh fruit, fresh potatoes while lower consumption of alcoholic drinks compared to the other UK countries.

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Of all these plots really only the `pairs()` plot was useful. This however took a bit of work to interpret and will not scale when I am looking at much bigger data sets.

## PCA the rescue

The main function in “base R” for PCA is called `prcomp()`.

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.7e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.0e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.0e+00


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```

# Create a data frame for plotting
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

# Plot PC1 vs PC2 with ggplot
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()

```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.


```

# Create a data frame for plotting
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

my_cols <- c("Wales"="red", "England"="orange", "Scotland"="blue", "N.Ireland"="darkgreen")

```

```
# Plot PC1 vs PC2 with ggplot
ggplot(df, aes(x = PC1, y = PC2, label = Country, color = Country)) +
  geom_point(size = 3, col="grey") +
  geom_text(vjust = -0.5) +
  scale_color_manual(values = my_cols) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.699876e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00


Q. How much variance is captured in the first PC? 67.4%

Q. How many PCs do I need to capture at least 90% of the total variance in the dataset? Two PCs capture 96.5% of the total variance.

Q. Plot our main PCA result. Folks can call this different things depending on their field of study e.g. “PC plot”, “ordination plot”, “Score plot”, “PC1 vs. PC2 plot”...

```
# Create scree plot with ggplot
variance_df <- data.frame(
  PC = factor(paste0("PC", 1:length(v)), levels = paste0("PC", 1:length(v))),
  Variance = v
)

ggplot(variance_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  xlab("Principal Component") +
  ylab("Percent Variation") +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 0))
```



```
attributes(pca)

$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"


$class
[1] "prcomp"
```

To generate our PCA score plot we want the `pca$x` component of the result object.

```
pca$x
```


	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	1.612425e-14
Wales	-240.52915	-224.646925	-56.475555	4.751043e-13
Scotland	-91.86934	286.081786	-44.415495	-6.044349e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.145386e-13

```
plot(pca$x[,1], pca$x[,2])
```




```
library(ggplot2)

ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col=my_cols)
```




```
my_cols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=my_cols, pch=16)
```



### Digging deeper (variable loadings)


How do the original variables (i.e., the 17 different foods) contribute to our new PCs?

```
## Lets focus on PC1 as it accounts for > 90% of variance
ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about? Soft drinks and alcoholic drinks.

```
ggplot(pca$rotation) +
  aes(x = PC2,
      y = reorder(rownames(pca$rotation), PC2)) +
  geom_col(fill = "lightgreen") +
  xlab("PC2 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



## PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

```
nrow(rna.data)
```

```
[1] 100
```

```
ncol(rna.data)
```


```
[1] 10
```

Q10: How many genes and samples are in this data set? How many PCs do you think it will take to have a useful overview of this data set (see below)? 100 genes and 10 samples. Based on the PC1 vs PC2 plot, the first two PCs already seem to separate the two samples meaningfully.

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

# Create data frame for plotting
df <- as.data.frame(pca$x)
df$Sample <- rownames(df)

## Plot with ggplot
ggplot(df) +
  aes(x = PC1, y = PC2, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, size = 3) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251
	PC8	PC9	PC10				
Standard deviation	0.62065	0.60342	3.39e-15				
Proportion of Variance	0.00385	0.00364	0.00e+00				
Cumulative Proportion	0.99636	1.00000	1.00e+00				


```
# Calculate variance explained
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)

# Create scree plot data
scree_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  Variance = pca.var[1:10]
)
```

```

ggplot(scree_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  ggtitle("Quick scree plot") +
  xlab("Principal Component") +
  ylab("Variance") +
  theme_bw()

```



```

## Percent variance is often more informative to look at
pca.var.per

```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```

# Create percent variance scree plot
scree_pct_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  PercentVariation = pca.var.per[1:10]
)


ggplot(scree_pct_df) +
  aes(x = PC, y = PercentVariation) +

```

```

geom_col(fill = "steelblue") +
ggtitle("Scree Plot") +
xlab("Principal Component") +
ylab("Percent Variation") +
theme_bw()

```



```

## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

# Add condition to data frame
df$condition <- substr(df$Sample, 1, 2)
df$color <- colvec

ggplot(df) +
  aes(x = PC1, y = PC2, color = color, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, hjust = 0.5, show.legend = FALSE) +
  scale_color_identity() +
  xlab(paste0("PC1 (", pca.var.per[1], "%)")) +
  ylab(paste0("PC2 (", pca.var.per[2], "%)"))

```

```
theme_bw()
```

