

Class 08: Analysis of Breast Cancer Cells

Amy Nguyen (PID: A18148284)

Table of contents

Background	1
Data Import	1
Principal Component Analysis (PCA)	5
Hierachial Clustering	14
Combining Methods	15
Sensitivity/Specificity	17
Prediction	18

Background

In today's class we will be employing all the R techniques for data analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

Data Import

The data is in CSV format:

```
fna.data <- "WisconsinCancer.csv"  
  
wisc.df <- data.frame(fna.data, row.names=1)  
  
wisc.df <- read.csv("WisconsinCancer.csv", row.names = 1)
```

wee peak at the data

```

head(wisc.df, 3)

  diagnosis radius_mean texture_mean perimeter_mean area_mean
842302      M     17.99     10.38     122.8     1001
842517      M     20.57     17.77     132.9     1326
84300903    M     19.69     21.25     130.0     1203
  smoothness_mean compactness_mean concavity_mean concave.points_mean
842302      0.11840     0.27760     0.3001     0.14710
842517      0.08474     0.07864     0.0869     0.07017
84300903    0.10960     0.15990     0.1974     0.12790
  symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
842302      0.2419      0.07871     1.0950     0.9053     8.589
842517      0.1812      0.05667     0.5435     0.7339     3.398
84300903    0.2069      0.05999     0.7456     0.7869     4.585
  area_se smoothness_se compactness_se concavity_se concave.points_se
842302      153.40     0.006399     0.04904     0.05373     0.01587
842517      74.08      0.005225     0.01308     0.01860     0.01340
84300903    94.03      0.006150     0.04006     0.03832     0.02058
  symmetry_se fractal_dimension_se radius_worst texture_worst
842302      0.03003     0.006193     25.38      17.33
842517      0.01389     0.003532     24.99      23.41
84300903    0.02250     0.004571     23.57      25.53
  perimeter_worst area_worst smoothness_worst compactness_worst
842302      184.6       2019      0.1622      0.6656
842517      158.8       1956      0.1238      0.1866
84300903    152.5       1709      0.1444      0.4245
  concavity_worst concave.points_worst symmetry_worst
842302      0.7119      0.2654      0.4601
842517      0.2416      0.1860      0.2750
84300903    0.4504      0.2430      0.3613
  fractal_dimension_worst
842302      0.11890
842517      0.08902
84300903    0.08758

```

Q1. How many observations are in this dataset? 569 observations.

```
nrow(wisc.df)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis? 212 observations.

```
sum( wisc.df$diagnosis == "M" )
```

```
[1] 212
```

```
table( wisc.df$diagnosis)
```

```
B     M  
357 212
```

Q3. How many variables/features in the data are suffixed with _mean? 10 variables

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867

```

84358402 94.44      0.011490      0.02461      0.05688      0.01885
843786   27.19      0.007510      0.03345      0.03672      0.01137
          symmetry_se fractal_dimension_se radius_worst texture_worst
842302      0.03003      0.006193      25.38       17.33
842517      0.01389      0.003532      24.99       23.41
84300903    0.02250      0.004571      23.57       25.53
84348301    0.05963      0.009208      14.91       26.50
84358402    0.01756      0.005115      22.54       16.67
843786      0.02165      0.005082      15.47       23.75
          perimeter_worst area_worst smoothness_worst compactness_worst
842302        184.60      2019.0       0.1622      0.6656
842517        158.80      1956.0       0.1238      0.1866
84300903    152.50      1709.0       0.1444      0.4245
84348301      98.87      567.7       0.2098      0.8663
84358402      152.20      1575.0       0.1374      0.2050
843786        103.40      741.6       0.1791      0.5249
          concavity_worst concave.points_worst symmetry_worst
842302        0.7119      0.2654       0.4601
842517        0.2416      0.1860       0.2750
84300903    0.4504      0.2430       0.3613
84348301      0.6869      0.2575       0.6638
84358402      0.4000      0.1625       0.2364
843786        0.5355      0.1741       0.3985
          fractal_dimension_worst
842302          0.11890
842517          0.08902
84300903        0.08758
84348301        0.17300
84358402        0.07678
843786          0.12440

```

```
length(grep("_mean", colnames(wisc.df), value=T))
```

```
[1] 10
```

```
colnames(wisc.df)
```

```
[1] "diagnosis"                  "radius_mean"
[3] "texture_mean"                "perimeter_mean"
[5] "area_mean"                   "smoothness_mean"
[7] "compactness_mean"             "concavity_mean"
```

```
[9] "concave.points_mean"      "symmetry_mean"
[11] "fractal_dimension_mean"   "radius_se"
[13] "texture_se"              "perimeter_se"
[15] "area_se"                 "smoothness_se"
[17] "compactness_se"          "concavity_se"
[19] "concave.points_se"        "symmetry_se"
[21] "fractal_dimension_se"    "radius_worst"
[23] "texture_worst"           "perimeter_worst"
[25] "area_worst"               "smoothness_worst"
[27] "compactness_worst"        "concavity_worst"
[29] "concave.points_worst"     "symmetry_worst"
[31] "fractal_dimension_worst"
```

We need to remove the `diagnosis` column before we do any further analysis of this data set - we don't want to pass this to PCA etc. We will save it as a separate wee vector that we can use later to compare our findings to those experts.

```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df$diagnosis
```

Principal Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features/dimensions are on very different scales in the original data set.

```
colMeans(wisc.data)
```

<code>radius_mean</code>	<code>texture_mean</code>	<code>perimeter_mean</code>
1.412729e+01	1.928965e+01	9.196903e+01
<code>area_mean</code>	<code>smoothness_mean</code>	<code>compactness_mean</code>
6.548891e+02	9.636028e-02	1.043410e-01
<code>concavity_mean</code>	<code>concave.points_mean</code>	<code>symmetry_mean</code>
8.879932e-02	4.891915e-02	1.811619e-01
<code>fractal_dimension_mean</code>	<code>radius_se</code>	<code>texture_se</code>
6.279761e-02	4.051721e-01	1.216853e+00
<code>perimeter_se</code>	<code>area_se</code>	<code>smoothness_se</code>
2.866059e+00	4.033708e+01	7.040979e-03
<code>compactness_se</code>	<code>concavity_se</code>	<code>concave.points_se</code>
2.547814e-02	3.189372e-02	1.179614e-02
<code>symmetry_se</code>	<code>fractal_dimension_se</code>	<code>radius_worst</code>

```

2.054230e-02      3.794904e-03      1.626919e+01
texture_worst      perimeter_worst      area_worst
2.567722e+01      1.072612e+02      8.805831e+02
smoothness_worst   compactness_worst   concavity_worst
1.323686e-01      2.542650e-01      2.721885e-01
concave.points_worst symmetry_worst fractal_dimension_worst
1.146062e-01      2.900756e-01      8.394582e-02

```

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale=T)
```

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14

```

Standard deviation      0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
Cumulative Proportion  0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
                           PC15    PC16    PC17    PC18    PC19    PC20    PC21
Standard deviation      0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
Cumulative Proportion  0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
                           PC22    PC23    PC24    PC25    PC26    PC27    PC28
Standard deviation      0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
Cumulative Proportion  0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
                           PC29    PC30
Standard deviation      0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion  1.00000 1.00000

```

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)? 44.27%

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 3 PCs

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? 7 PCs

```
attributes(wisc.pr)
```

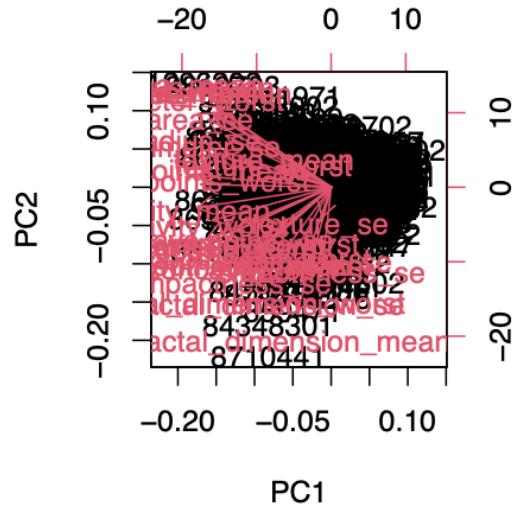
```

$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"

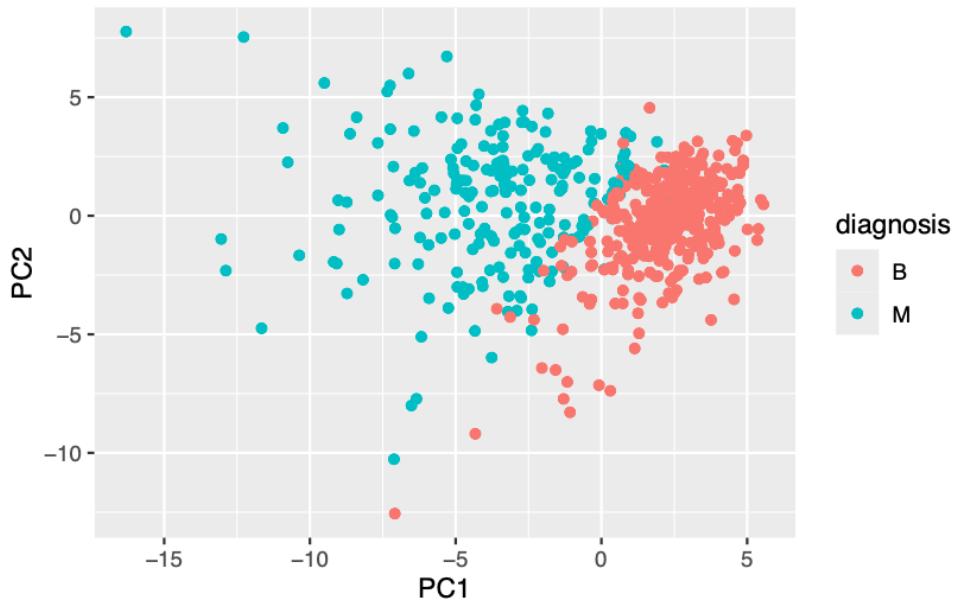
```

```
biplot(wisc.pr)
```



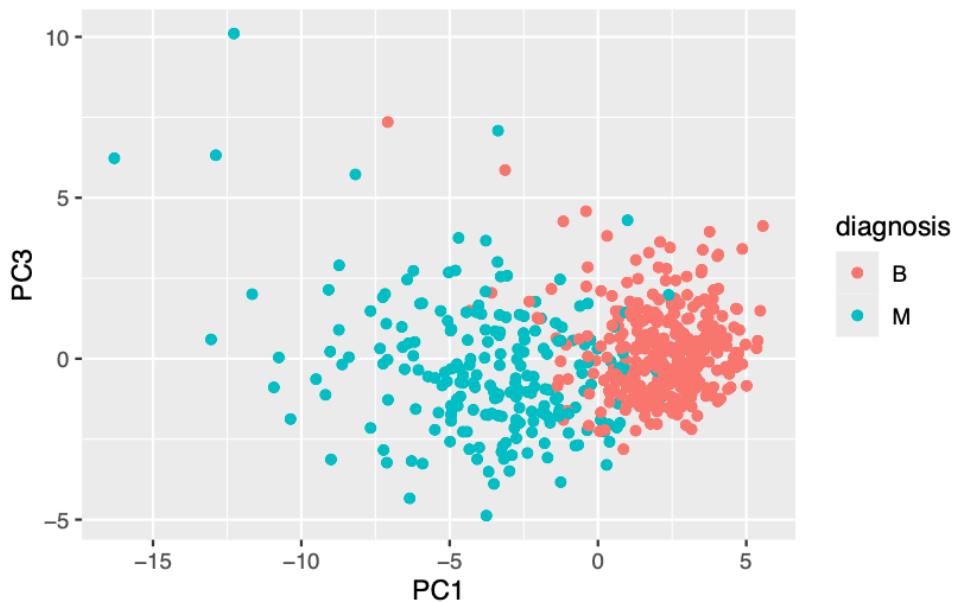
Q7. What stands out to you about this plot? Is it easy or difficult to understand?
 Why? The plot is extremely poor for visual interpretation and difficult to understand because the row names are plotted which obsecures the data points.

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots? These plots show how PC1 successfully captures the visual distinction between the benign (red data points) and malignant (blue data points) cancer samples.

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



Calculate variance of each component

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

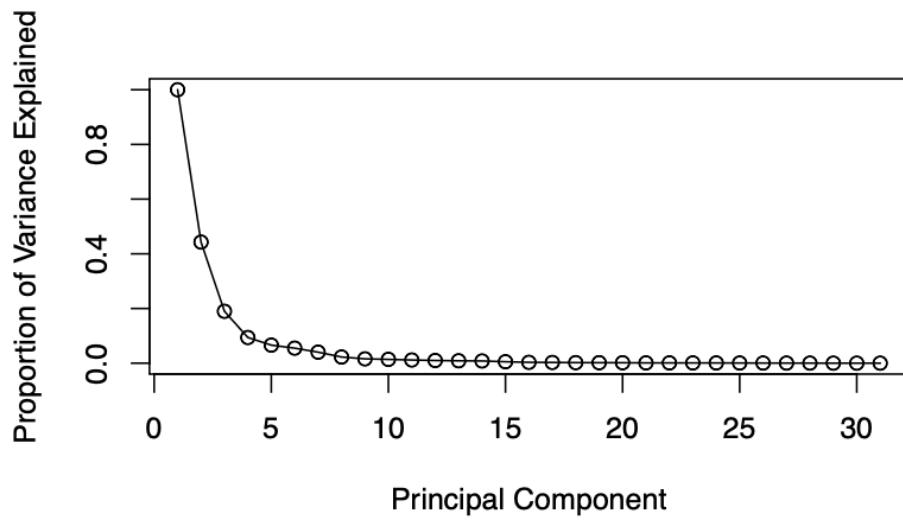
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

Variance explained by each principal component: pve

```
pve <- pr.var / sum(pr.var)
```

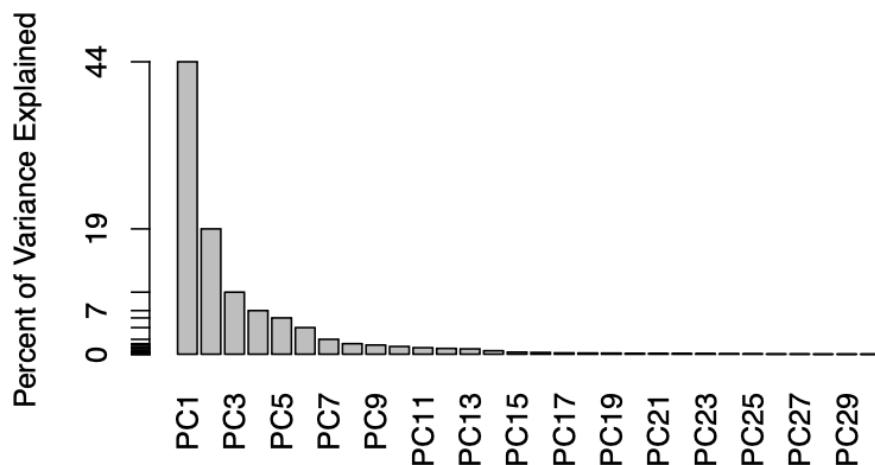
Plot variance explained for each principal component

```
plot(c(1,pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



Alternative scree plot of the same data, note data driven y-axis

```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

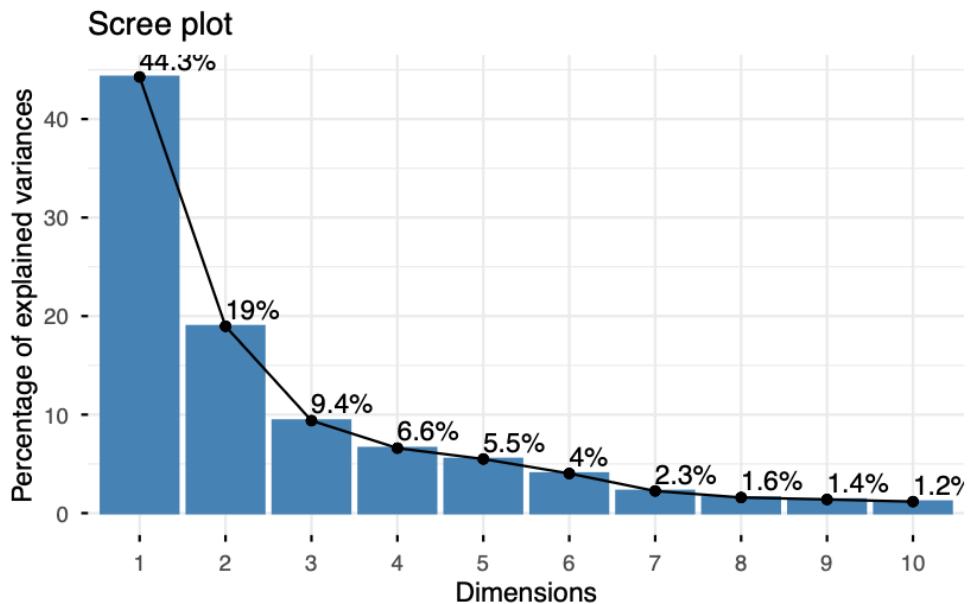


```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



```
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

```
pc1_contributions <- wisc.pr$rotation[, 1]
abs(pc1_contributions) > abs(pc1_contributions["concave.points_mean"])
```

radius_mean	texture_mean	perimeter_mean
FALSE	FALSE	FALSE
area_mean	smoothness_mean	compactness_mean
FALSE	FALSE	FALSE
concavity_mean	concave.points_mean	symmetry_mean
FALSE	FALSE	FALSE
fractal_dimension_mean	radius_se	texture_se
FALSE	FALSE	FALSE
perimeter_se	area_se	smoothness_se
FALSE	FALSE	FALSE
compactness_se	concavity_se	concave.points_se
FALSE	FALSE	FALSE
symmetry_se	fractal_dimension_se	radius_worst
FALSE	FALSE	FALSE

```

texture_worst      perimeter_worst      area_worst
      FALSE            FALSE            FALSE
smoothness_worst   compactness_worst   concavity_worst
      FALSE            FALSE            FALSE
concave.points_worst symmetry_worst fractal_dimension_worst
      FALSE            FALSE            FALSE

```

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one? The component of the loading vector for the feature `concave.points_mean` is -0.26. There are no features with larger contributions.

Hierarchical Clustering

The goal of this section is to do hierarchical clustering of the original data to see if there is any obvious grouping into malignant and benign clusters.

In short, these results are not good!

First we will scale our `wisc.data`

```

data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method = "complete")

```

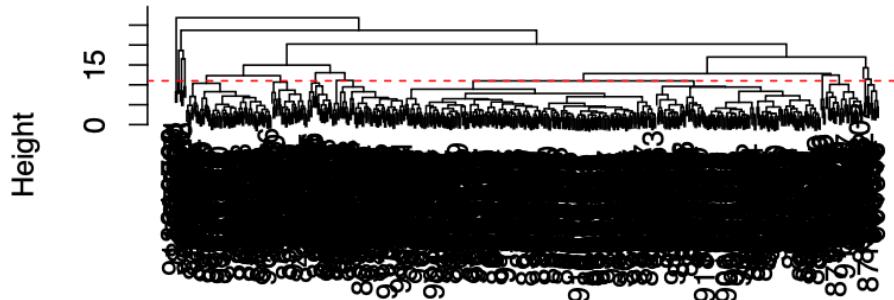
Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters? At height = 11

```

plot(wisc.hclust)
abline(h = 11, col="red", lty=2)

```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)  
table(wisc.hclust.clusters, diagnosis)
```

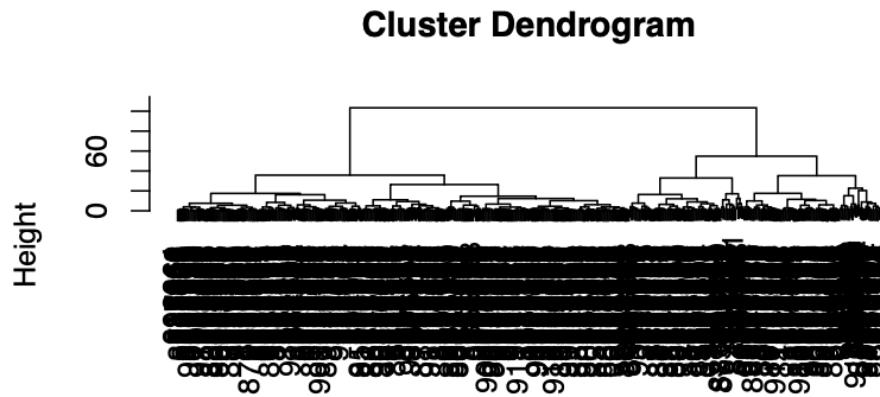
```
diagnosis  
wisc.hclust.clusters   B    M  
1      12 165  
2       2   5  
3     343  40  
4       0   2
```

Q12. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning. The ward.D2 method gives my favorite result. It produces the most well-separated clusters as it minimizes the within-cluster variance.

Combining Methods

The idea here is that I can take my new variables (i.e. the scores on the PCs `wisc.pr$x`) that are better descriptors of the data-set than the original features (i.e. the 30 columns in `wisc.data`) and use these as a basis for the clustering.

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
plot(wisc.pr.hclust)
```



pc.dist
hclust (*, "ward.D2")

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters)
```

```
wisc.pr.hclust.clusters
 1  2
203 366
```

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

Q13. How well does the newly created hclust model with two clusters separate out the two "M" and "B" diagnoses?

I can now run `table()` with both my clustering `wisc.pr.hclust.clusters` and the expert diagnoses

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	24	179
2	333	33

Our cluster “1” has 179 “M” diagnosis Our cluster “2” has 333 “B” diagnosis

179 TP 24 FP 333 TN 33 FN

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.hclust.clusters` and `wisc.pr.hclust.clusters`) with the vector containing the actual diagnoses. The `wisc.hclust.clusters` model did worse at separating diagnoses compared to the `wisc.pr.hclust.clusters` model.

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	24	179
2	333	33

Sensitivity/Specificity

Sensitivity: $TP / (TP + FN)$

179/(179+33)

[1] 0.8443396

Specificity: TN/(TN+FP)

333/(333+24)

[1] 0.9327731

Q15. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity? The model based on PCA-transformed data resulted in the best specificity (93%) and sensitivity (84%).

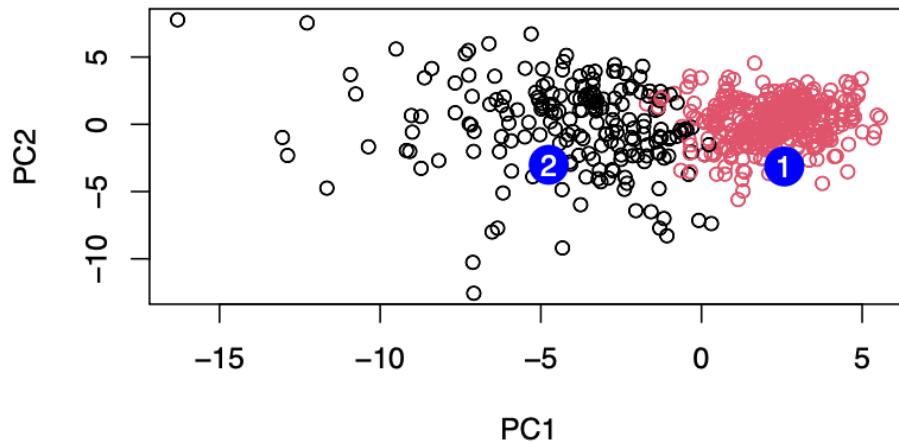
Prediction

We can use our PCA model for prediction of new unseen cases.

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=wisc.pr$hclust.clusters)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results? We should prioritize patient 2 for follow up.